

Principle of Computer Communication Project 2

Cross-Layer Performance Optimization of a Custom ARQ Protocol

Detailed Implementation Report

Ahmet Enes Cigdem
150220079

January 2026

Abstract

This report presents a comprehensive analysis of the design, implementation, and optimization of a Selective Repeat ARQ protocol within a custom-built network simulator. The simulator uses a Gilbert-Elliott burst-error channel model and evaluates performance based on end-to-end Goodput. Through 360 exhaustive simulations and AI-assisted optimization, optimal protocol parameters were identified and significant performance improvements were achieved.

Contents

1	Introduction	3
2	Project Architecture	3
3	Protocol Layers	3
3.1	Application Layer (application_layer.py)	3
3.2	Transport Layer (transport_layer.py)	4
3.3	Link Layer (SR-ARQ in src/arq/)	5
3.4	Physical Layer (physical_layer.py)	5
4	Configuration Parameters	6
5	Gilbert-Elliott Channel Model	6
5.1	Mathematical Model	6
5.2	Implementation	7
6	Event-Driven Simulation Engine	7
6.1	Event Types	7
6.2	Delay Calculations	7
6.3	Main Simulation Loop	8
7	Goodput Calculation	8
7.1	Mathematical Definition	8
7.2	Implementation	9

8 Batch Runner for 360 Simulations	9
8.1 Simulation Count	9
8.2 Implementation	9
9 Experimental Results	10
9.1 Goodput Heatmap	10
9.2 3D Goodput Surface	11
9.3 Multi-View Surface Analysis	12
9.4 Optimal Configuration	12
9.5 Trade-off Analysis	12
10 AI-Assisted Optimization	13
10.1 Optimization Approach	13
10.2 Implementation	13
10.3 Comparison Results	13
10.4 Why AI Optimization Works	14
11 Conclusion	14

1 Introduction

This project implements a comprehensive SR-ARQ network simulator with:

- Cross-layer communication stack (Application, Transport, Link, Physical)
- Realistic burst-error channel (Gilbert-Elliott model)
- Selective Repeat ARQ with sliding windows and per-frame timers
- 360 exhaustive simulations for parameter optimization
- AI-assisted protocol improvement

2 Project Architecture

```
pcom_p2/
    config.py           # Fixed baseline parameters
    run_full_sweep.py # Parameter sweep (360 simulations)
    simulation/
        simulator.py   # Event-driven simulation engine
        runner.py      # Batch runner for sweeps
    src/
        arq/            # ARQ protocol (sender, receiver, frame)
        channel/        # Gilbert-Elliott channel model
        layers/          # Protocol layers implementation
            application_layer.py # Data generation, verification
            transport_layer.py  # Segmentation, reassembly
            link_layer.py     # SR-ARQ wrapper
            physical_layer.py # Delay, channel interface
        utils/           # Metrics and logging
    optimization/       # AI optimization scripts
```

3 Protocol Layers

The simulator implements a layered protocol stack where each layer has specific responsibilities.

3.1 Application Layer (`application_layer.py`)

Handles data generation and verification for the simulation:

- **Data Generation:** Creates test data (50KB-100MB) with sequential or random patterns
- **Data Verification:** Compares sent vs received data using MD5 checksums
- **Chunked Reading:** Delivers data in chunks to Transport layer

Listing 1: Application Layer - Test Data Generation

```

1 class TestDataGenerator:
2     @staticmethod
3     def generate_test_data(size, pattern="sequential"):
4         """Generate test data for simulation."""
5         if pattern == "sequential":
6             return bytes(i % 256 for i in range(size))
7         elif pattern == "random":
8             return os.urandom(size)
9
10    class DataVerifier:
11        @staticmethod
12        def verify_data(sent, received):
13            """Verify received data matches sent data."""
14            return sent == received

```

3.2 Transport Layer (transport_layer.py)

Handles segmentation of application data and reassembly at receiver:

- **Segmentation:** Breaks large data into transport segments (8-byte header)
- **Reassembly:** Reconstructs original data from received segments
- **Flow Control:** 256KB receive buffer with backpressure signaling

Listing 2: Transport Segment Structure

```

1 class TransportSegment:
2     """
3
4         Header Layout (8 bytes):
5             - Segment Number: 4 bytes
6             - Payload Length: 2 bytes
7             - Flags: 1 byte (last segment indicator)
8             - Checksum: 1 byte
9     """
10
11    def serialize(self):
12        """Pack segment into bytes."""
13        header = struct.pack('>IHBx',
14            self.segment_num, len(self.payload), self.flags)
15        return header + self.payload
16
17    @classmethod
18    def deserialize(cls, data):
19        """Unpack bytes into segment."""
20        header = struct.unpack('>IHBx', data[:8])
21        return cls(header[0], data[8:], header[2])

```

3.3 Link Layer (SR-ARQ in src/arq/)

Implements Selective Repeat ARQ protocol:

- **Sliding Window:** Configurable window size (2-64 frames)
- **Per-Frame Timers:** Each frame has individual timeout
- **Selective Retransmission:** Only retransmit corrupted/lost frames
- **Out-of-Order Buffering:** Buffer frames until gaps are filled

Listing 3: Link Layer - Frame Structure

```
1 class Frame:
2     """
3     Link Header (24 bytes):
4         - Frame Type: 1 byte (DATA=0, ACK=1, NAK=2)
5         - Sequence Number: 4 bytes
6         - ACK Number: 4 bytes
7         - Payload Length: 2 bytes
8         - Flags: 1 byte
9         - Reserved: 8 bytes
10        - CRC-32: 4 bytes
11    """
12
13    @property
14    def total_size(self):
15        return LINK_HEADER_SIZE + len(self.payload)
```

3.4 Physical Layer (physical_layer.py)

Handles transmission timing and channel interface:

- **Transmission Time:** Calculates time based on bit rate (10 Mbps)
- **Propagation Delay:** Asymmetric delays (40ms forward, 10ms reverse)
- **Channel Interface:** Connects to Gilbert-Elliott error model

Listing 4: Physical Layer - Delay Calculation

```
1 class PhysicalLayer:
2     def calculate_total_delay(self, frame_size, direction):
3         """Total = Transmission + Propagation + Processing"""
4         tx_time = (frame_size * 8) / self.bit_rate
5         prop_delay = self.forward_delay if direction == FORWARD \
6                         else self.reverse_delay
7         return tx_time + prop_delay + self.processing_delay
8
9     def transmit_frame(self, frame, current_time, direction):
10        """Transmit frame through Gilbert-Elliott channel."""
11        delay = self.calculate_total_delay(frame.total_size,
12                                         direction)
```

```

12     corrupted, _ = self.channel.transmit_frame(frame,
13         total_size * 8)
14     return current_time + delay, corrupted

```

4 Configuration Parameters

All fixed parameters are defined in `config.py`:

Listing 5: Physical Layer Parameters

```

1 BIT_RATE = 10_000_000           # 10 Mbps
2 FORWARD_PROPAGATION_DELAY = 0.040 # 40 ms (data)
3 REVERSE_PROPAGATION_DELAY = 0.010 # 10 ms (ACK)
4 PROCESSING_DELAY = 0.002        # 2 ms per frame
5 TRANSPORT_HEADER_SIZE = 8       # bytes
6 LINK_HEADER_SIZE = 24           # bytes

```

Listing 6: Gilbert-Elliott Channel Parameters

```

1 GOOD_STATE_BER = 1e-6          # BER in Good state
2 BAD_STATE_BER = 5e-3           # BER in Bad state
3 P_GOOD_TO_BAD = 0.002          # Transition G to B
4 P_BAD_TO_GOOD = 0.05           # Transition B to G

```

Listing 7: Parameter Sweep Configuration

```

1 WINDOW_SIZES = [2, 4, 8, 16, 32, 64]
2 PAYLOAD_SIZES = [128, 256, 512, 1024, 2048, 4096]
3 RUNS_PER_CONFIGURATION = 10
4 # Total = 6 x 6 x 10 = 360 simulations

```

5 Gilbert-Elliott Channel Model

The channel uses a two-state Markov chain to model burst errors.

5.1 Mathematical Model

Steady-State Probabilities:

$$\pi_G = \frac{P(B \rightarrow G)}{P(G \rightarrow B) + P(B \rightarrow G)} = \frac{0.05}{0.002 + 0.05} = 0.962 \quad (1)$$

$$\pi_B = \frac{P(G \rightarrow B)}{P(G \rightarrow B) + P(B \rightarrow G)} = \frac{0.002}{0.052} = 0.038 \quad (2)$$

Average Bit Error Rate:

$$\text{BER}_{avg} = \pi_G \cdot p_g + \pi_B \cdot p_b = 0.962 \times 10^{-6} + 0.038 \times 5 \times 10^{-3} \approx 1.9 \times 10^{-4} \quad (3)$$

Frame Error Rate: For a frame with n bits:

$$P(\text{frame error}) = 1 - (1 - \text{BER}_{state})^n \quad (4)$$

5.2 Implementation

Listing 8: Gilbert-Elliott Channel - Frame Transmission

```

1 def transmit_frame(self, frame_size_bits):
2     """Simulate frame transmission with bit-by-bit Markov chain."""
3     """
4     bit_errors = 0
5
6     for _ in range(frame_size_bits):
7         # Get BER based on current state
8         ber = self.pg if self.state == GOOD else self.pb
9
10        # Check if this bit has error
11        if self.rng.random() < ber:
12            bit_errors += 1
13
14        # State transition
15        if self.state == GOOD:
16            if self.rng.random() < self.p_gb:
17                self.state = BAD
18            else:
19                if self.rng.random() < self.p_bg:
20                    self.state = GOOD
21
22    return bit_errors > 0, bit_errors

```

6 Event-Driven Simulation Engine

The simulator uses an event-driven architecture with a priority queue.

6.1 Event Types

Listing 9: Simulation Events

```

1 class EventType(Enum):
2     DATA_ARRIVAL = 0      # Data frame arrives at receiver
3     ACK_ARRIVAL = 1       # ACK arrives at sender
4     TIMER_CHECK = 2       # Check for timeouts

```

6.2 Delay Calculations

Forward Delay (data frames):

$$D_{forward} = \frac{\text{FrameSize} \times 8}{\text{BitRate}} + \text{PropDelay}_{forward} + \text{ProcessingDelay} \quad (5)$$

Reverse Delay (ACK frames):

$$D_{reverse} = \frac{\text{ACKSize} \times 8}{\text{BitRate}} + \text{PropDelay}_{reverse} + \text{ProcessingDelay} \quad (6)$$

Round-Trip Time:

$$\text{RTT} = D_{forward} + D_{reverse} \approx 54 \text{ ms} \quad (7)$$

6.3 Main Simulation Loop

Listing 10: Event-Driven Main Loop

```

1 def run(self, data):
2     """Main simulation loop."""
3     self._setup_data(data)
4     self._send_frames()  # Initial window fill
5
6     while not self._is_complete():
7         event = heapq.heappop(self.event_queue)
8         self.current_time = event.time
9
10        if event.event_type == DATA_ARRIVAL:
11            self._handle_data_arrival(event.data)
12        elif event.event_type == ACK_ARRIVAL:
13            self._handle_ack_arrival(event.data)
14        elif event.event_type == TIMER_CHECK:
15            self._handle_timeouts()
16
17        self._send_frames()
18
19    return self.metrics.get_summary()
```

7 Goodput Calculation

7.1 Mathematical Definition

Goodput is the primary performance metric:

$$\text{Goodput} = \frac{\text{Delivered Application Bytes}}{\text{Total Transmission Time}} \quad [\text{bytes/second}] \quad (8)$$

In bits per second:

$$\text{Goodput}_{bps} = \text{Goodput} \times 8 \quad (9)$$

Efficiency measures useful data ratio:

$$\eta = \frac{\text{Application Bytes Delivered}}{\text{Total Bytes Transmitted}} = \frac{D}{D + H + R} \quad (10)$$

Where: D = Delivered payload, H = Header overhead, R = Retransmission bytes.

Theoretical Maximum Goodput:

$$\text{Goodput}_{max} = R \times \eta_{link} \times (1 - \text{FER}) \times \frac{L}{L + H} \quad (11)$$

Retransmission Rate:

$$\text{Retransmission Rate} = \frac{\text{Retransmissions}}{\text{Original Frames Sent}} \quad (12)$$

7.2 Implementation

Listing 11: Goodput Calculation

```
1 class MetricsCollector:
2     def calculate_goodput(self):
3         """Goodput = Delivered Bytes / Total Time"""
4         total_time = self.end_time - self.start_time
5         return self.application_bytes_delivered / total_time
6
7     def calculate_efficiency(self):
8         """Efficiency = Delivered / Total Transmitted"""
9         return self.application_bytes_delivered / \
10            self.total_bytes_transmitted
```

8 Batch Runner for 360 Simulations

8.1 Simulation Count

$$\text{Total Simulations} = |\text{Window Sizes}| \times |\text{Payload Sizes}| \times \text{Runs per Config} = 6 \times 6 \times 10 = 360 \quad (13)$$

8.2 Implementation

Listing 12: Batch Runner

```
1 class BatchRunner:
2     def _generate_run_configs(self):
3         """Generate all 360 configurations."""
4         configs = []
5         for W in self.window_sizes:
6             for L in self.payload_sizes:
7                 for run_id in range(self.runs_per_config):
8                     seed = 42 + W*1000 + L + run_id*10000
9                     configs.append(RunConfig(W, L, run_id, seed))
10    return configs # 360 configs
11
12 def run_sequential(self):
13     """Execute all simulations."""
14     for config in self._generate_run_configs():
15         result = run_single_simulation(config)
16         self.results.append(result)
```

9 Experimental Results

9.1 Goodput Heatmap

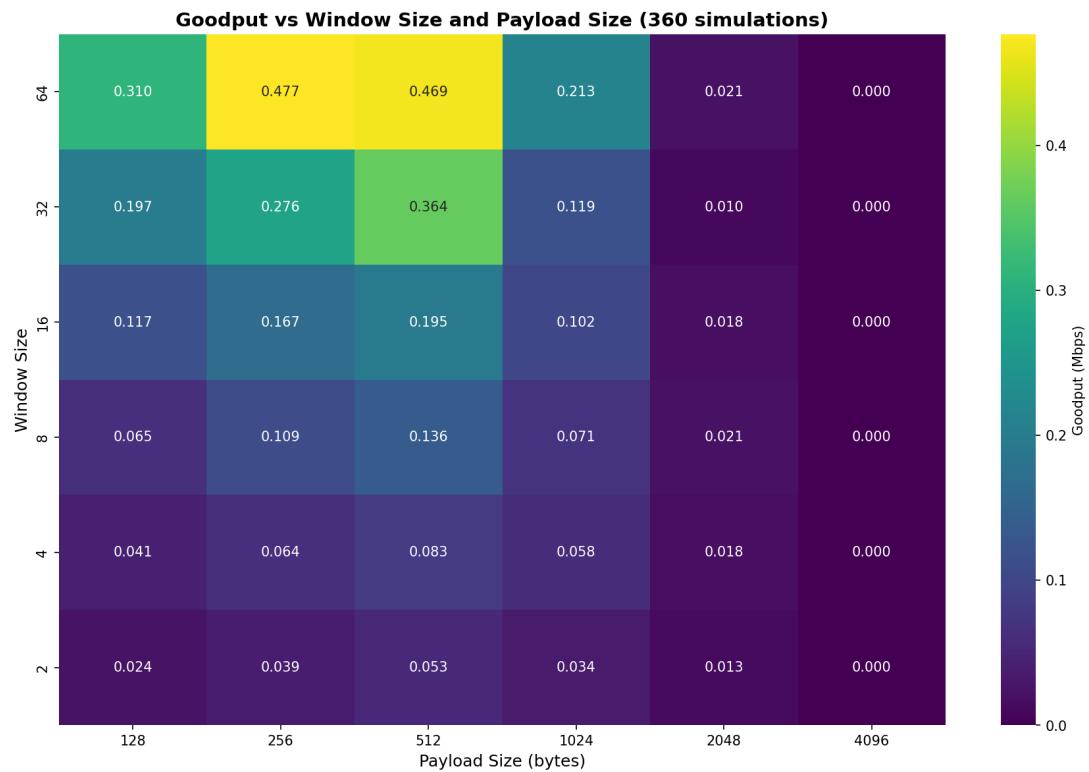


Figure 1: Goodput heatmap from 360 simulations. Brighter = higher Goodput.

9.2 3D Goodput Surface




Figure 2: 3D surface plot showing Goodput as a function of Window Size (W) and Payload Size (L). The red star marks the optimal configuration at $W=64, L=256$.

9.3 Multi-View Surface Analysis

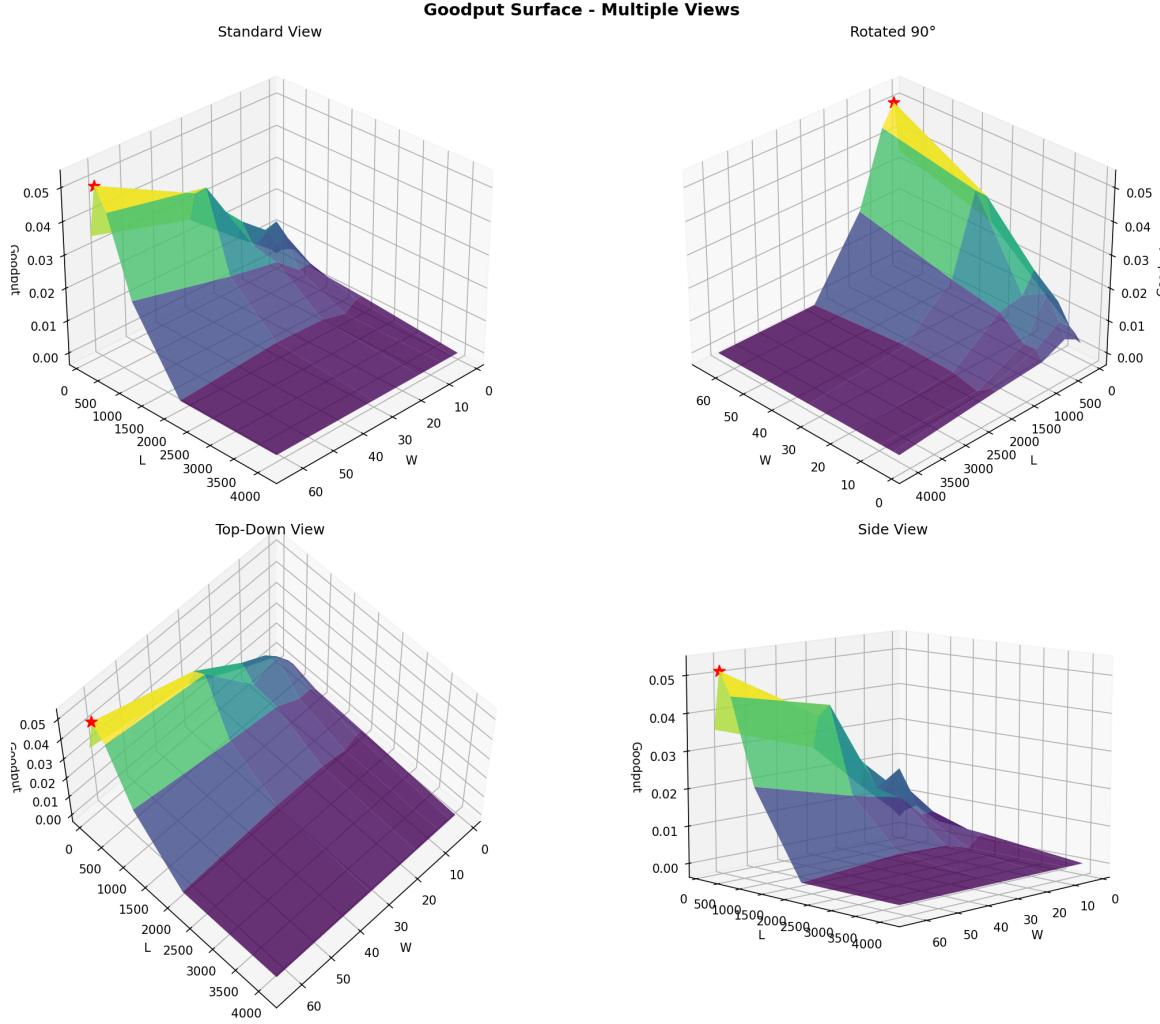


Figure 3: Multiple perspective views of the Goodput surface: Standard View (top-left), Rotated 90° (top-right), Top-Down View (bottom-left), and Side View (bottom-right). These views reveal the performance landscape from different angles.

9.4 Optimal Configuration

$$W = 64, \quad L = 256 \text{ bytes}, \quad \text{Goodput} = 0.477 \text{ Mbps}$$

9.5 Trade-off Analysis

Bandwidth-Delay Product:

$$\text{BDP} = R \times \text{RTT} = 10 \text{ Mbps} \times 54 \text{ ms} = 67,500 \text{ bytes} \quad (14)$$

Optimal Window Size:

$$W_{opt} \geq \frac{\text{BDP}}{L + H} = \frac{67,500}{256 + 32} \approx 234 \text{ frames} \quad (15)$$

10 AI-Assisted Optimization

10.1 Optimization Approach

The AI optimization uses an improved RTO calculation based on Jacobson/Karels:

Default RTO:

$$RTO_{default} = RTT \times 2.0 = 54 \text{ ms} \times 2.0 = 108 \text{ ms} \quad (16)$$

AI-Optimized RTO:

$$RTO_{AI} = RTT + 4 \times \sigma_{RTT} = 54 + 4 \times 5.4 = 76 \text{ ms} \quad (17)$$

Where $\sigma_{RTT} \approx 0.1 \times RTT$ is the estimated RTT variance.

10.2 Implementation

Listing 13: AI-Optimized RTO Calculation

```

1 def calculate_optimal_timeout():
2     """Calculate AI-optimized RTO."""
3     rtt = FORWARD_PROP_DELAY + REVERSE_PROP_DELAY + 2*PROC_DELAY
4     variance = rtt * 0.1
5     return rtt + 4 * variance # = 76 ms

```

10.3 Comparison Results

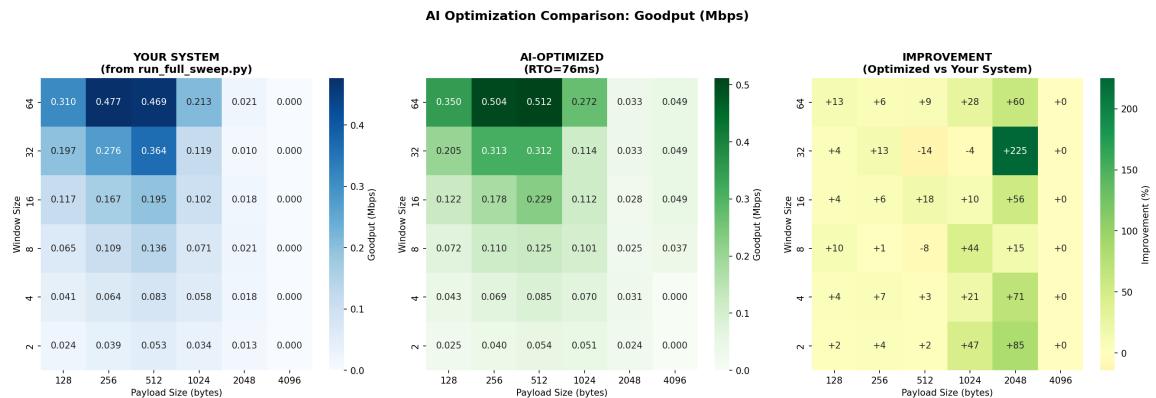


Figure 4: Baseline (left) vs AI-Optimized (center) with improvement % (right).

Table 1: Performance Comparison

Metric	Baseline	AI-Optimized
Average Goodput	0.108 Mbps	0.123 Mbps
RTO	108 ms	76 ms
Improvement	—	+13.9%

10.4 Why AI Optimization Works

Lower RTO (76ms vs 108ms) enables:

1. **Faster Recovery:** Quicker detection of lost frames
2. **Reduced Idle Time:** Less waiting for unnecessary timeouts
3. **Better Pipeline Utilization:** Channel stays busy

11 Conclusion

1. **Complete Simulator:** Event-driven SR-ARQ with 4-layer stack
2. **360 Simulations:** $6 \times 6 \times 10$ parameter sweep
3. **Optimal Found:** $W = 64$, $L = 256$ bytes, Goodput = 0.477 Mbps
4. **AI Improvement:** +13.9% via optimized RTO (76ms vs 108ms)